

CSE 4074 – Programming Assignment Report

Socket Programming – HTTP Server and Proxy Server

Project Members:

Cem Eren Kula 150120059

Doğukan Onmaz 150120071

Ferhat Sirkeci 150120067

HTTP Server

Overview

The `HttpServer.py` file implements a simple HTTP server capable of handling client requests. It leverages multi-threading to handle multiple connections concurrently. The server responds to GET requests, checks the validity of the request URI, and generates appropriate HTTP responses.

Features

1. Multi-threaded Design:

- Each client connection is handled in a separate thread, ensuring the server can process multiple requests simultaneously.

2. Request Parsing:

- Validates the HTTP method (supports GET only).
- Checks the URI for integer values within a specified range (100 to 20,000).
- Handles non-integer URIs (e.g., `favicon.ico`).

3. Response Generation:

- Constructs HTTP responses based on the status of the request (200 OK, 400 Bad Request, or 501 Not Implemented).
- Allows for optional saving of responses to HTML files in a "Responses" directory.

4. Port Configurability:

- The port on which the server listens can be specified via command-line arguments.

Code Details

1. Request Parsing:

- The `parse_request()` function validates the request method and URI.
- Unsupported methods return a 501 Not Implemented status, while invalid URIs return a 400 Bad Request.

2. Response Generation:

- The `generate_response()` function creates HTML responses based on the status code.
- If enabled, it saves responses as files in the Responses directory.

3. **Concurrency:**

- The `handle_client()` function uses threading to ensure independent handling of client connections.
- A shared counter (`response_counter`) tracks the number of responses generated, protected by a thread-safe lock.

4. **Server Setup:**

- The `start_server()` function initializes a socket and listens for incoming connections.
- Client connections are accepted and handled in separate threads.

Potential Enhancements

We can add a cache mechanism to store frequently requested responses, reducing processing time for repeated requests.

Proxy Server

Overview

ProxyServer.py implements a simple proxy server designed to forward HTTP requests to a specified web server. The proxy enforces restrictions on allowed hosts and ports, ensuring controlled access to backend resources. It employs multi-threading to handle multiple client connections concurrently.

Features

1. **Request Handling:**

- Parses incoming HTTP requests, extracting the method, URI, and HTTP version.
- Handles HTTP requests with the `http://` scheme, routing them to the appropriate destination.
- Verifies host and port against predefined allowed values.

2. **Request Forwarding:**

- Rewrites the request line and headers before forwarding them to the target web server.
- Establishes a connection to the web server and sends the modified request.

3. **Response Transmission:**

- Receives responses from the web server and relays them back to the client without modification.

4. **Access Control:**

- Implements strict access controls by allowing requests only to specified hosts and ports (localhost:8080 in this implementation).
- Returns an HTTP 403 Forbidden status for disallowed requests.

5. **Multi-threading:**

- Uses threading to handle multiple client connections simultaneously, ensuring high responsiveness.

Code Details

1. **Request Parsing:**

- Extracts and validates the method, URI, and HTTP version.
- Parses the host and port from the URI for routing.

2. **Access Control:**

- The `handle_client()` function checks if the requested host and port match the allowed configuration (`ALLOWED_HOST` and `ALLOWED_PORT`).
- Disallowed requests return a 403 Forbidden response.

3. **Request Forwarding:**

- Constructs a modified HTTP request with updated headers.
- Forwards the request to the specified backend server using a new socket connection.

4. **Response Handling:**

- Relays responses from the backend server back to the client.
- Returns a 502 Bad Gateway response if the backend server is unreachable.

5. **Thread Management:**

- Each client connection is managed in a separate thread using Python's threading module.
- Threads operate independently to ensure concurrent handling of requests.

Potential Enhancements

We can also add caching to proxy server.

Stress Testing Proxy Server

Note: In first 6 test we are only changing Number of Threads.

1. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 200

Result:
  total: 6134 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 302.4236ms
  latency median: 310.19295ms
  latency average: 320.501337ms
  latency max: 894.4106ms
  transfers: 3067000 bytes
  rps: 624.0 requests per sec
```

2. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 10 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 10
  duration: 10s
  connections: 200

Result:
  total: 6110 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 302.2514ms
  latency median: 310.70325ms
  latency average: 322.224815ms
  latency max: 920.3025ms
  transfers: 3055000 bytes
  rps: 620.7 requests per sec
```

3. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 15 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 15
  duration: 10s
  connections: 200

Result:
  total: 6134 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 302.0881ms
  latency median: 311.08475ms
  latency average: 321.976489ms
  latency max: 900.845ms
  transfers: 3067000 bytes
  rps: 621.2 requests per sec
```

4. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 20 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 20
  duration: 10s
  connections: 200

Result:
  total: 6118 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 302.476ms
  latency median: 309.9598ms
  latency average: 321.733574ms
  latency max: 899.4672ms
  transfers: 3059000 bytes
  rps: 621.6 requests per sec
```

5. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 30 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 30
  duration: 10s
  connections: 200

Result:
  total: 4052 requests
  errors: 200 errors
  error percentage: 4.9%
  latency min: 302.4908ms
  latency median: 310.4739ms
  latency average: 410.429161ms
  latency max: 2.0464625s
  transfers: 1926000 bytes
  rps: 487.3 requests per sec
```

6. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 50 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 50
  duration: 10s
  connections: 200

Result:
  total: 6104 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 302.0708ms
  latency median: 311.0276ms
  latency average: 322.28509ms
  latency max: 903.5281ms
  transfers: 3052000 bytes
  rps: 620.6 requests per sec
```

Note: Now we are going to change parallel connection number.

7. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 300
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 300

Result:
  total: 8954 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 302.1653ms
  latency median: 312.5088ms
  latency average: 329.316436ms
  latency max: 1.4126371s
  transfers: 4477000 bytes
  rps: 911.0 requests per sec
```

8. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 400
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 400

Result:
  total: 7604 requests
  errors: 400 errors
  error percentage: 5.3%
  latency min: 302.5775ms
  latency median: 334.75815ms
  latency average: 461.589804ms
  latency max: 2.1780513s
  transfers: 3602000 bytes
  rps: 866.6 requests per sec
```


9. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 500
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 500

Result:
  total: 2046 requests
  errors: 1984 errors
  error percentage: 97.0%
  latency min: 318.199ms
  latency median: 2.03017325s
  latency average: 2.011188254s
  latency max: 2.2217041s
  transfers: 31000 bytes
  rps: 248.6 requests per sec
```

10. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 600
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 600

Result:
  total: 9363 requests
  errors: 272 errors
  error percentage: 2.9%
  latency min: 302.1415ms
  latency median: 368.1708ms
  latency average: 581.849947ms
  latency max: 3.49527s
  transfers: 4545632 bytes
  rps: 1031.2 requests per sec
```

11. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 700
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 700

Result:
  total: 8104 requests
  errors: 620 errors
  error percentage: 7.7%
  latency min: 302.5394ms
  latency median: 374.95625ms
  latency average: 774.456714ms
  latency max: 3.5605587s
  transfers: 3743045 bytes
  rps: 903.9 requests per sec
```

12. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 2000
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 2000

Result:
  total: 10259 requests
  errors: 3089 errors
  error percentage: 30.1%
  latency min: 303.3555ms
  latency median: 1.5844329s
  latency average: 1.606819021s
  latency max: 4.8227412s
  transfers: 3609981 bytes
  rps: 1244.7 requests per sec
```

Note: Now we are going to change parallel connection number and thread number at the same time.

13. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 30 -d 10 -c 2000
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 30
  duration: 10s
  connections: 2000

Result:
  total: 2485 requests
  errors: 2411 errors
  error percentage: 97.0%
  latency min: 438.5654ms
  latency median: 2.6913847s
  latency average: 2.873178881s
  latency max: 6.8132719s
  transfers: 37000 bytes
  rps: 696.1 requests per sec
```

Stress Testing HTTP Server

Note: We are going to change only Number of Threads.

1. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 200

Result:
  total: 6217 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 301.3448ms
  latency median: 308.8747ms
  latency average: 318.573079ms
  latency max: 887.6678ms
  transfers: 3108500 bytes
  rps: 627.8 requests per sec
```

2. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 10 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 10
  duration: 10s
  connections: 200

Result:
  total: 6174 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 301.4203ms
  latency median: 309.4557ms
  latency average: 319.792877ms
  latency max: 896.8709ms
  transfers: 3087000 bytes
  rps: 625.4 requests per sec
```

3. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 15 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 15
  duration: 10s
  connections: 200

Result:
  total: 6124 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 301.4323ms
  latency median: 310.2879ms
  latency average: 320.593177ms
  latency max: 891.4693ms
  transfers: 3062000 bytes
  rps: 623.8 requests per sec
```

4. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 20 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 20
  duration: 10s
  connections: 200

Result:
  total: 3981 requests
  errors: 200 errors
  error percentage: 5.0%
  latency min: 301.4705ms
  latency median: 309.9849ms
  latency average: 410.639546ms
  latency max: 2.0723956s
  transfers: 1890500 bytes
  rps: 487.0 requests per sec
```

5. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 30 -d 10 -c 200
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 30
  duration: 10s
  connections: 200

Result:
  total: 6173 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 301.3014ms
  latency median: 309.0182ms
  latency average: 319.443004ms
  latency max: 895.8791ms
  transfers: 3086500 bytes
  rps: 626.1 requests per sec
```

Note: Now we are going to change only parallel connection number.

6. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 300
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 300

Result:
  total: 8991 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 301.3568ms
  latency median: 311.3676ms
  latency average: 327.527089ms
  latency max: 1.3682045s
  transfers: 4495500 bytes
  rps: 916.0 requests per sec
```

7. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 400
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 400

Result:
  total: 7622 requests
  errors: 539 errors
  error percentage: 7.1%
  latency min: 301.4365ms
  latency median: 314.3253ms
  latency average: 476.914463ms
  latency max: 2.3016451s
  transfers: 3541500 bytes
  rps: 838.7 requests per sec
```

8. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 5 -d 10 -c 500
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 5
  duration: 10s
  connections: 500

Result:
  total: 2074 requests
  errors: 2000 errors
  error percentage: 96.4%
  latency min: 310.3929ms
  latency median: 2.0356221s
  latency average: 2.019142679s
  latency max: 2.3006094s
  transfers: 37000 bytes
  rps: 247.6 requests per sec
```

Note: Now we are going to also increase thread number with 500 connections.

9. Test

```
C:\Users\cemer>winrk http://localhost:8080/500 -t 20 -d 10 -c 500
Input:
  url: http://localhost:8080/500
  method: GET
  threads: 20
  duration: 10s
  connections: 500

Result:
  total: 13553 requests
  errors: 0 errors
  error percentage: 0.0%
  latency min: 301.2999ms
  latency median: 339.9516ms
  latency average: 362.492534ms
  latency max: 1.4065992s
  transfers: 6776500 bytes
  rps: 1379.3 requests per sec
```

Conclusion

When conducting stress tests on the proxy server, the following behaviors were observed:

1. Impact of Increasing Parallel Connections:

- As the number of parallel connections increases, both the error rate and average latency rise significantly.
- This occurs because the server struggles to manage a higher volume of simultaneous requests with a fixed number of threads.

2. Effect of Increasing Thread Count:

- By increasing the number of threads, the server can handle a greater degree of concurrency more efficiently.
- Higher thread counts allow the server to respond to simultaneous requests with reduced latency and error rates.

3. Thread Count vs. Latency and Errors:

- Simply increasing the thread count without increasing the number of parallel connections does not affect latency or error rates.
- This demonstrates that the server's performance is primarily limited by the demand for concurrency rather than the thread pool size under low-load conditions.

These insights highlight the importance of tuning thread count relative to expected traffic patterns for optimal performance during high-concurrency scenarios.